



PARALLEL AND DISTRIBUTED COMPUTING

An introduction



SHARED MEMORY PROGRAMMING

ROADMAP

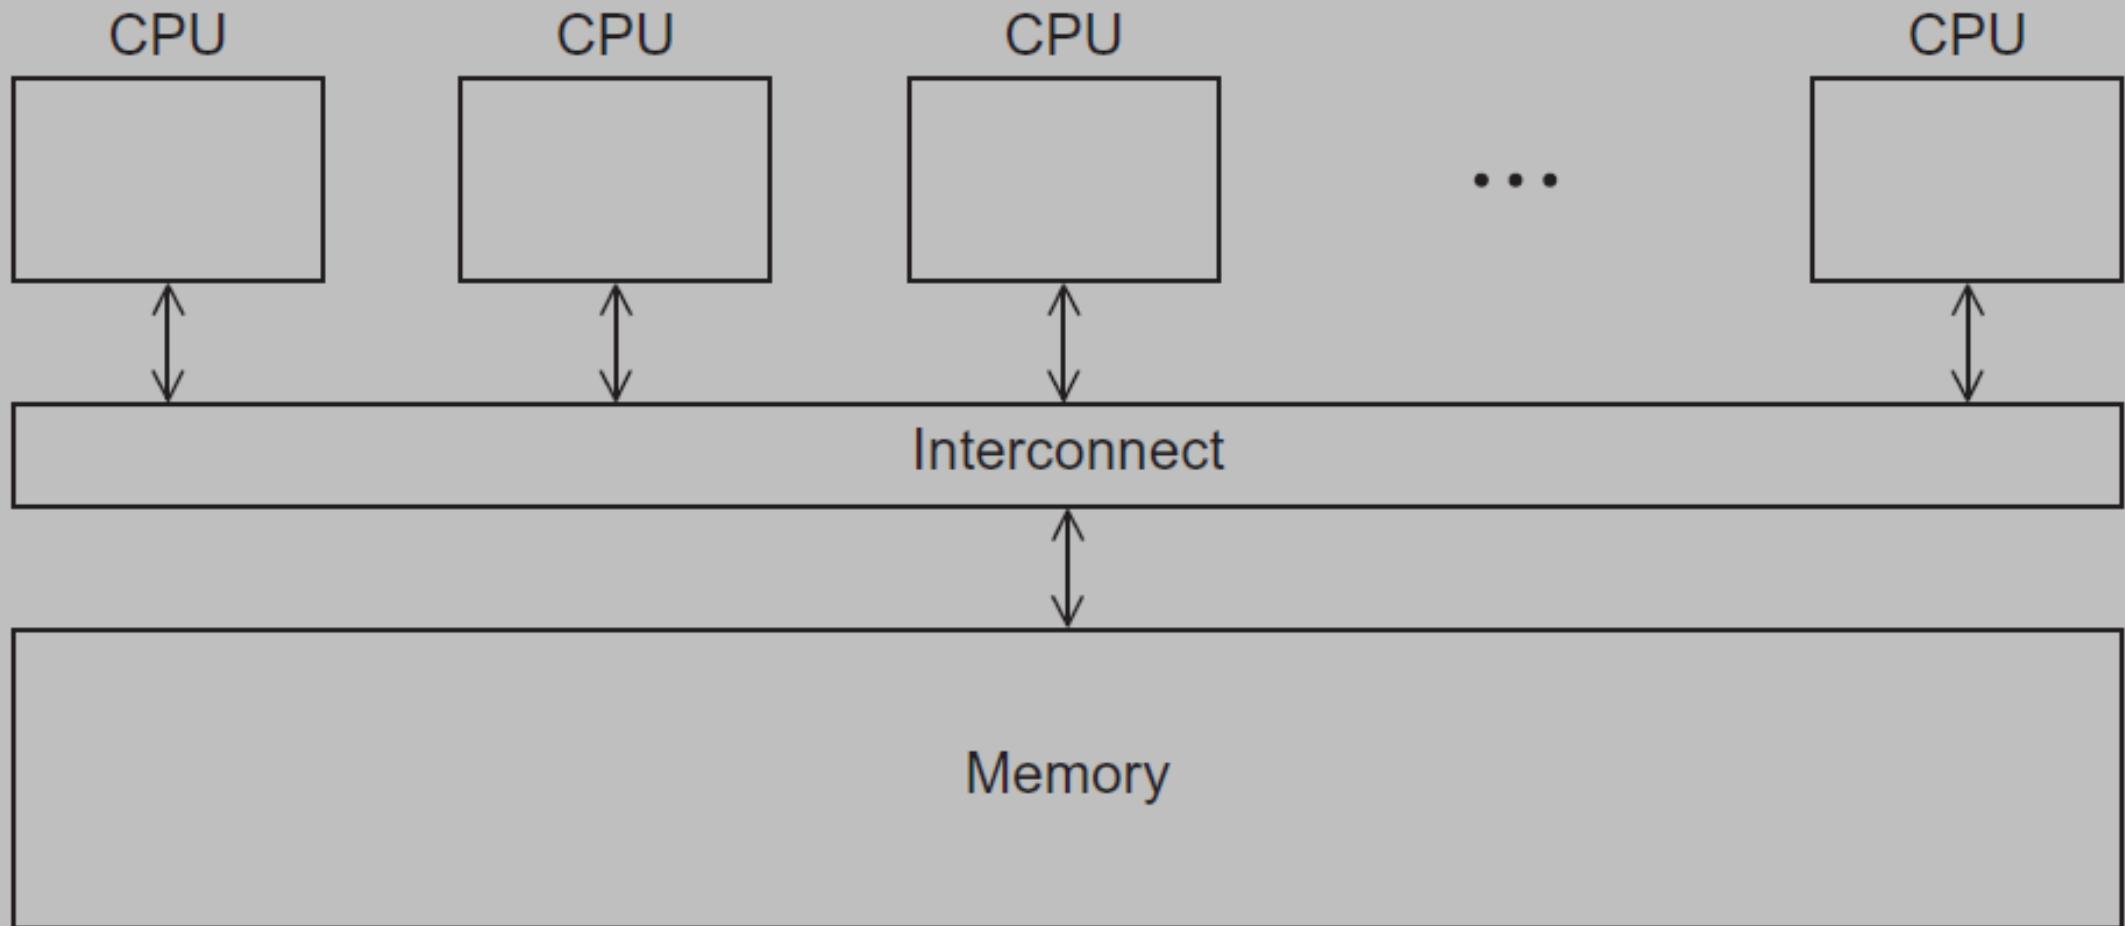
- Writing programs that use OpenMP.
- Using OpenMP to parallelize many serial for loops with only small changes to the source code.
- Task parallelism.
- Explicit thread synchronization.
- Standard problems in shared-memory programming.

OPENMP



- An API for shared-memory parallel programming.
- MP = multiprocessing
- Designed for systems in which each thread or process can potentially have access to all available memory.
- System is viewed as a collection of cores or CPU's, all of which have access to main memory.

A SHARED MEMORY SYSTEM





- Special preprocessor instructions.

#pragma

- Typically added to a system to allow behaviors that aren't part of the basic C specification.
- Compilers that don't support the pragmas ignore them.



```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
        Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);

} /* Hello */
```



```
gcc -g -Wall -fopenmp -o omp_hello omp_hello.c
```

```
./omp_hello 4
```

compiling

running with 4 threads

Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

possible outcomes

Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4

Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4

OPENMP PRAGMAS



- `# pragma omp parallel`
- Most basic parallel directive.
- The number of threads that run the following structured block of code is determined by the run-time system.

A PROCESS FORKING AND JOINING TWO THREADS





- Text that modifies a directive.
- The `num_threads` clause can be added to a parallel directive.
- It allows the programmer to specify the number of threads that should execute the following block.



pragma omp parallel num_threads (thread_count)



OF NOTE...

- There may be system-defined limitations on the number of threads that a program can start.
- The OpenMP standard doesn't guarantee that this will actually start thread_count threads.
- Most current systems can start hundreds or even thousands of threads.
- Unless we're trying to start a lot of threads, we will almost always get the desired number of threads.

SOME TERMINOLOGY

- In OpenMP parlance the collection of threads executing the *parallel block*
- the original thread and the new threads:
 - is called a team,
 - the original thread is called the master,
 - and the additional threads are called slaves.



IN CASE THE COMPILER DOESN'T SUPPORT OPENMP

```
# include <omp.h>
```

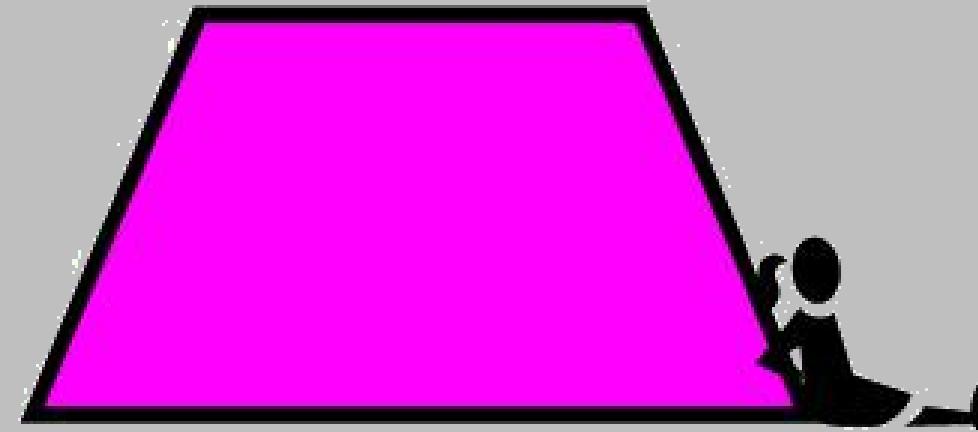


```
#ifdef _OPENMP  
# include <omp.h>  
#endif
```



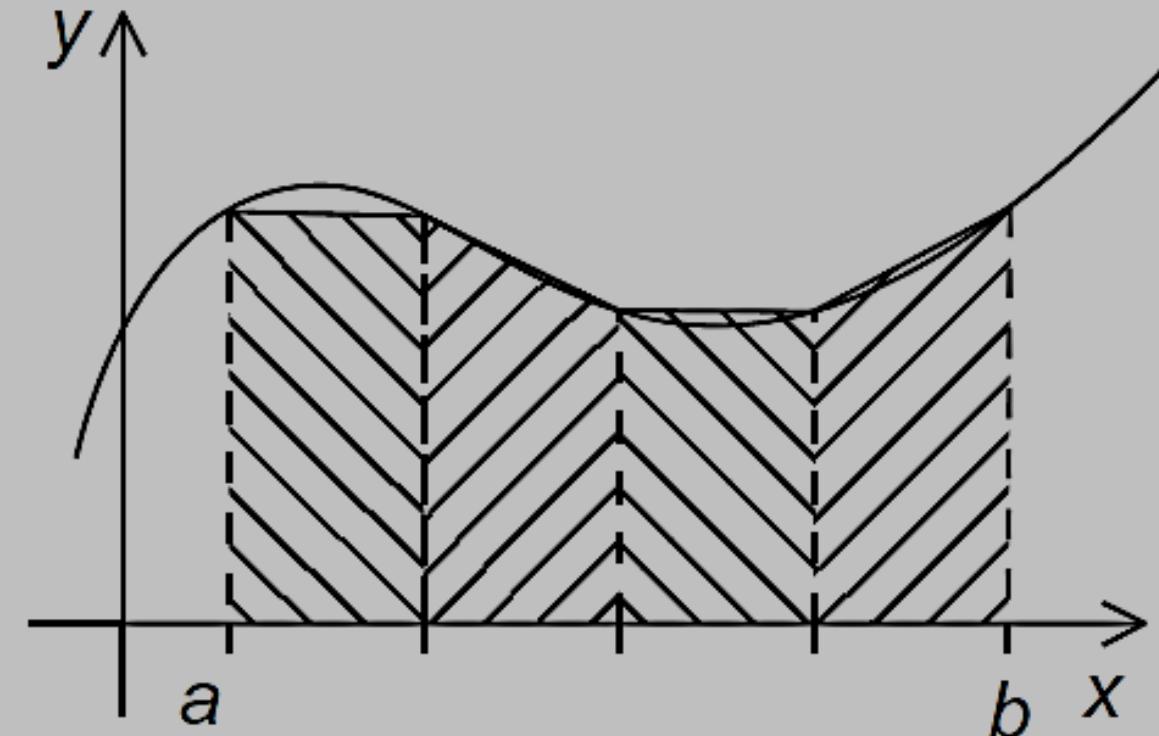
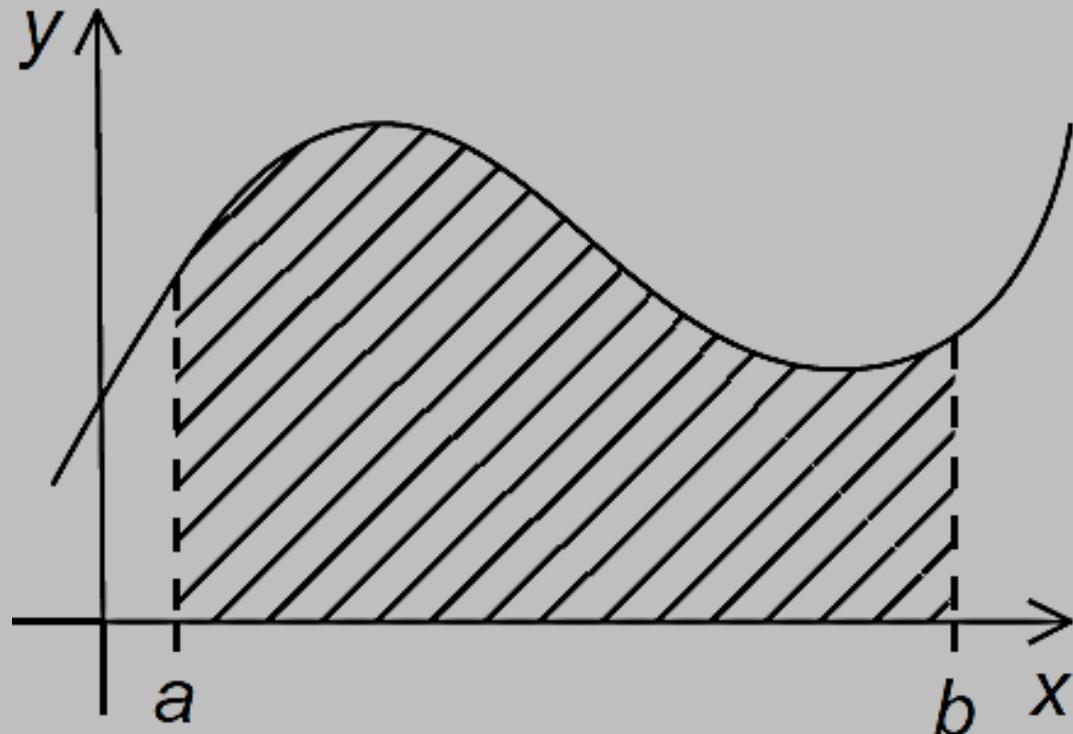
IN CASE THE COMPILER DOESN'T SUPPORT OPENMP

```
#ifdef _OPENMP  
    int my_rank =  
        omp_get_thread_num ();  
    int thread_count =  
        omp_get_num_threads ();  
  
#else  
    int my_rank = 0;  
    int thread_count = 1;  
  
#endif
```



THE TRAPEZOIDAL RULE

THE TRAPEZOIDAL RULE



SERIAL ALGORITHM

```
/* Input: a, b, n */
h = (b-a)/n;
approx = (f(a) + f(b))/2.0;
for (i = 1; i <= n-1; i++) {
    x_i = a + i*h;
    approx += f(x_i);
}
approx = h*approx;
```

A FIRST OPENMP VERSION

We identified two types of tasks:

1. Area Calculation

- Computation of the areas of individual trapezoids, and

2. Summation Of Areas

- Adding the areas of trapezoids

There is no communication among the tasks in the first collection, but each task in the first collection communicates with task 2.

A FIRST OPENMP VERSION

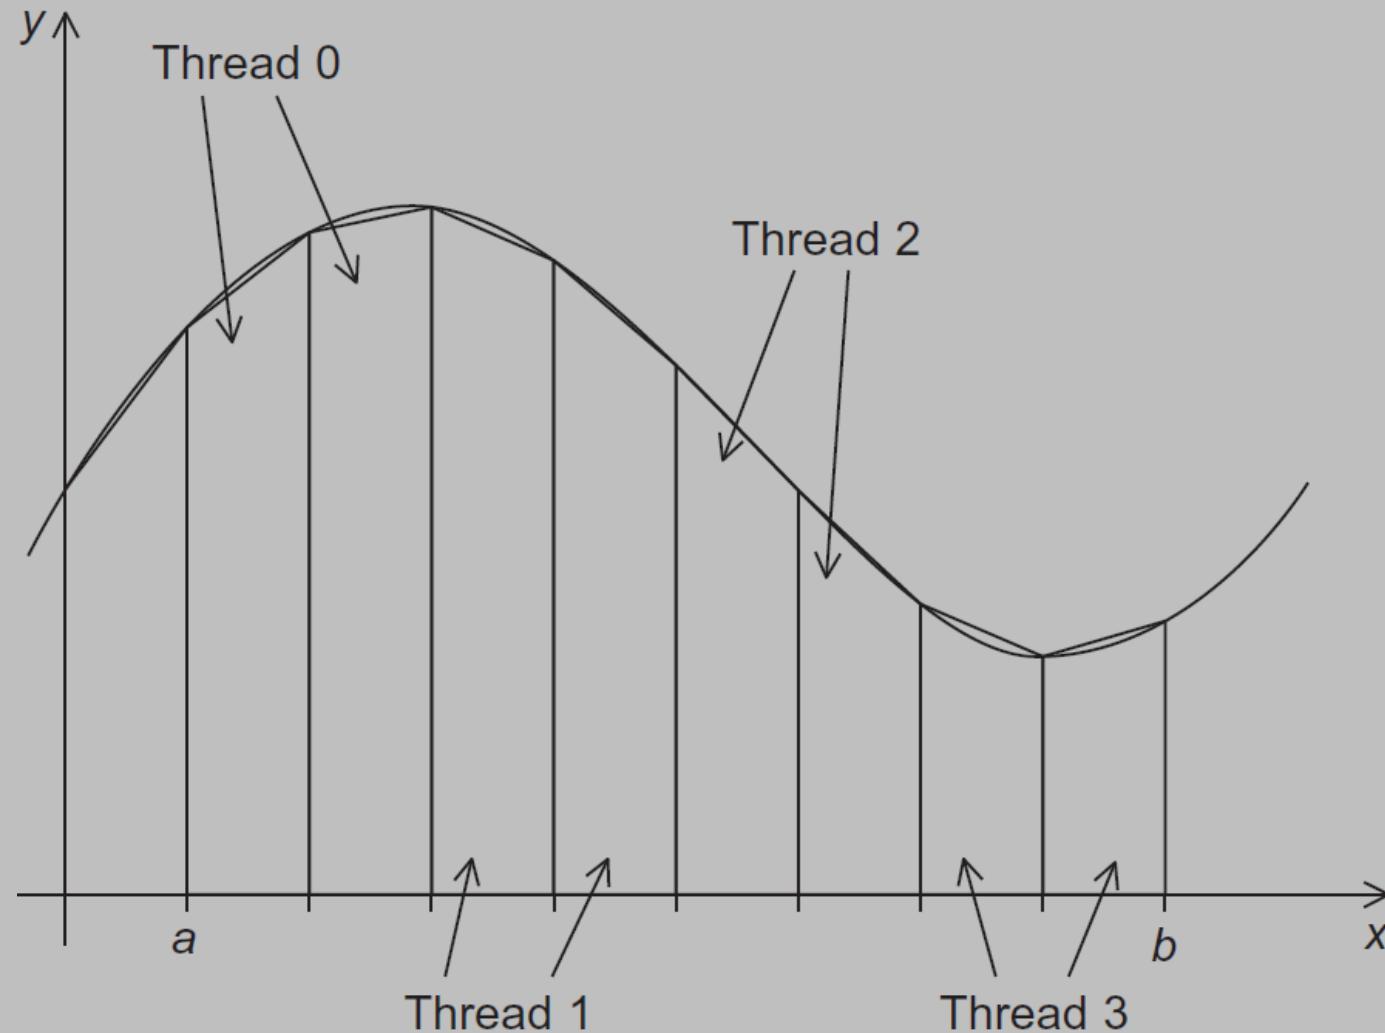


We assumed that there would be many more trapezoids than cores.



So, we aggregated tasks by assigning a contiguous block of trapezoids to each thread (and a single thread to each core).

ASSIGNMENT OF TRAPEZOIDS TO THREADS



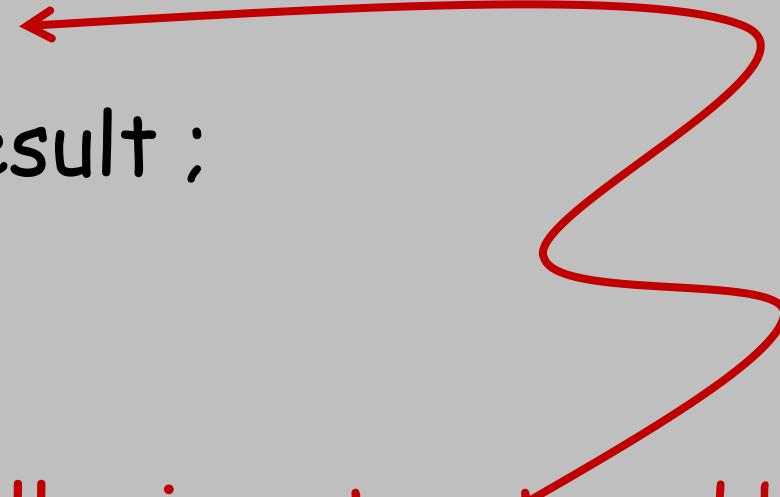
Time	Thread 0	Thread 1
0	global_result = 0 to register	finish my_result
1	my_result = 1 to register	global_result = 0 to register
2	add my_result to global_result	my_result = 2 to register
3	store global_result = 1	add my_result to global_result
4		store global_result = 2

Unpredictable results when two (or more) threads attempt to simultaneously execute:

`global_result += my_result ;`

MUTUAL EXCLUSION

```
# pragma omp critical  
global_result += my_result ;
```



only one thread can execute the following structured block
at a time



```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Trap(double a, double b, int n, double* global_result_p);

int main(int argc, char* argv[]) {
    double global_result = 0.0; /* Store result in global_result */
    double a, b; /* Left and right endpoints */
    int n; /* Total number of trapezoids */
    int thread_count;

    thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);
# pragma omp parallel num_threads(thread_count)
    Trap(a, b, n, &global_result);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n",
           a, b, global_result);
    return 0;
} /* main */
```

```
void Trap(double a, double b, int n, double* global_result_p) {
    double h, x, my_result;
    double local_a, local_b;
    int i, local_n;
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    h = (b-a)/n;
    local_n = n/thread_count;
    local_a = a + my_rank*local_n*h;
    local_b = local_a + local_n*h;
    my_result = (f(local_a) + f(local_b))/2.0;
    for (i = 1; i <= local_n-1; i++) {
        x = local_a + i*h;
        my_result += f(x);
    }
    my_result = my_result*h;

    # pragma omp critical
        *global_result_p += my_result;
} /* Trap */
```



SCOPE OF VARIABLES



- In serial programming, the scope of a variable consists of those parts of a program in which the variable can be used.
- In OpenMP, the scope of a variable refers to the set of threads that can access the variable in a parallel block.

SCOPE IN OPENMP

- A variable that can be accessed by all the threads in the team has shared scope.
- A variable that can only be accessed by a single thread has private scope.
- The default scope for variables declared before a parallel block is shared.





THE REDUCTION CLAUSE

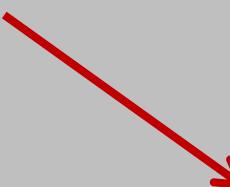
We need this more complex version to add each thread's local calculation to get *global_result*.

```
void Trap(double a, double b, int n, double* global_result_p);
```

Although we'd prefer this.

```
double Trap(double a, double b, int n);
```

```
global_result = Trap(a, b, n);
```



If we use this, there's no critical section!

```
double Local_trap(double a, double b, int n);
```

If we fix it like this...

```
global_result = 0.0;  
# pragma omp parallel num_threads(thread_count)  
{  
#     pragma omp critical  
        global_result += Local_trap(double a, double b, int n);  
}
```

... we force the threads to execute sequentially.

We can avoid this problem by declaring a private variable inside the parallel block and moving the critical section after the function call.

```
global_result = 0.0;  
# pragma omp parallel num_threads(thread_count)  
{  
    double my_result = 0.0; /* private */  
  
    my_result += Local_trap(double a, double b, int n);  
# pragma omp critical  
    global_result += my_result;  
}
```

X

“Do the
best you
can until
you know
better.”

Then when
you know
better, do
better.”

- Maya Angelou





REDUCTION OPERATORS

- A reduction operator is a binary operation (such as addition or multiplication).
- A reduction is a computation that repeatedly applies the same reduction operator to a sequence of operands in order to get a single result.
- All of the intermediate results of the operation should be stored in the same variable: the reduction variable.

A reduction clause can be added to a parallel directive.

```
reduction(<operator>: <variable list>)
```

+, *, -, &, |, ^, &&, ||

```
global_result = 0.0;  
# pragma omp parallel num_threads(thread_count) \  
    reduction(+: global_result)  
    global_result += Local_trap(double a, double b, int n);
```



THE "PARALLEL FOR" DIRECTIVE

PARALLEL FOR



- Forks a team of threads to execute the following structured block.
- However, the structured block following the parallel for directive must be a for loop.
- Furthermore, with the parallel for directive the system parallelizes the for loop by dividing the iterations of the loop among the threads.

```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```



```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+: approx)  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```

LEGAL FORMS FOR PARALLELIZABLE FOR STATEMENTS



```
for index = start ; index >= end ; index += incr
    index < end           index++
    index <= end          ++index
    index > end           index--
    index >= end          index--
    index < end           index--
    index <= end          --index
    index > end           index+=incr
    index >= end          index-=incr
    index < end           index+=incr
    index <= end          index+=incr
    index > end           index-=incr
```

CAVEATS

- The variable `index` must have integer or pointer type (e.g., it can't be a float).
- The expressions `start`, `end`, and `incr` must have a compatible type. For example, if `index` is a pointer, then `incr` must have integer type.



CAVEATS

- The expressions start, end, and incr must not change during execution of the loop.
- During execution of the loop, the variable index can only be modified by the "increment expression" in the for statement.

DATA DEPENDENCIES

```
fibo[ 0 ] = fibo[ 1 ] = 1;  
for (i = 2; i < n; i++)  
    fibo[ i ] = fibo[ i - 1 ] + fibo[ i - 2 ];  
  
fibo[ 0 ] = fibo[ 1 ] = 1;  
# pragma omp parallel for num_threads(2)  
for (i = 2; i < n; i++)  
    fibo[ i ] = fibo[ i - 1 ] + fibo[ i - 2 ];
```

↓

note 2 threads

↓

1 1 2 3 5 8 13 21 34 55

this is correct

↓

1 1 2 3 5 8 0 0 0 0

but sometimes
we get this

WHAT HAPPENED?

- OpenMP compilers don't check for dependences among iterations in a loop that's being parallelized with a parallel for directive.
- A loop in which the results of one or more iterations depend on other iterations cannot, in general, be correctly parallelized by OpenMP.