



Parallel Patterns

Array Sum, Trapezoidal Rule, Matrix Vector Multiplication

This Week....




- MPI – Array Sum
 - Reduction in MPI
- Trapezoidal Rule
- Matrix - Vector Multiplication



Array Sum in MPI using send and receive

The big picture





Sending Local size first then array


```
1.  for (i = 1; i < np - 1; ++i) {  
2.      index = i * local_size;  
3.      MPI_Send(&local_size, 1, MPI_INT, i, 0, MPI_COMM_WORLD);  
4.      MPI_Send(&arr[index], local_size, MPI_DOUBLE,  
5.                  i, 1, MPI_COMM_WORLD);  
6.  } //end for  
7. // last process adds remaining elements  
8. index = i * local_size;  
9. int elements_left = size - index;  
  
10. MPI_Send(&elements_left, 1, MPI_INT, i, 0, MPI_COMM_WORLD);  
11. MPI_Send(&arr[index], elements_left, MPI_DOUBLE,  
12.                  i, 1, MPI_COMM_WORLD);
```

Slave Code receiving local size then elements

```
1. else {//slave code
2.     printf("I am slave to He-Man!");
3.     MPI_Recv(&local_size, 1, MPI_INT,
4.             0, 0, MPI_COMM_WORLD, &status);
5.     printf(" He said %d!\n", local_size);
6.     // stores the received array segment in local array
arr2
7.     double* arr2 = malloc(sizeof(double) * local_size);
8.     MPI_Recv(arr2, local_size, MPI_DOUBLE,
9.             0, 1, MPI_COMM_WORLD, &status);
10.    printf(" Just Received my %d elements! Calculating
Local sum....\n", local_size);
```




Trapezoidal Rule




Trapezoidal rule in mpi

The Trapezoidal Rule



(a)



(b)

The Trapezoidal Rule


$$\text{Area of one trapezoid} = \frac{h}{2}[f(x_i) + f(x_{i+1})]$$

$$h = \frac{b - a}{n}$$

$$x_0 = a, x_1 = a + h, x_2 = a + 2h, \dots, x_{n-1} = a + (n-1)h, x_n = b$$

$$\text{Sum of trapezoid areas} = h[f(x_0)/2 + f(x_1) + f(x_2) + \cdots + f(x_{n-1}) + f(x_n)/2]$$

One trapezoid



Pseudo-code for a serial program

1. $h = (b-a)/n$
 - Height same for all
2. $estimate = (f(a) + f(b)) / 2.0$
 - 1st and last added once only
3. **for**($i = 1; i \leq n-1; i++$) {
 - Rest added twice
4. $x_i = a + i * h;$
 - Cancels divide by 2.0
5. $estimate += f(x_i);$
6. } //end for
7. $estimate = h * estimate$
 - Multiply by height.



- Partition problem solution into tasks.
- Identify communication channels between tasks.
- Aggregate tasks into composite tasks.
- Map composite tasks to cores.

Parallel Code 1

```
1.  Get_input(my_rank, comm_sz, &a, &b, &n);  
2.  h = (b - a) / n;  
3.  /* h is the same for all processes */  
4.  local_n = n / comm_sz;  
5.  /* So is the number of trapezoids */  
  
6.  /* Length of each process' interval of integration  
7.   * = local_n*h.  So my interval  
8.   * starts at: */  
9.  local_a = a + my_rank * local_n * h;  
10. local_b = local_a + local_n * h;  
11. local_int = Trap(local_a, local_b, local_n, h);
```

- Only process 0 gets from std_in, rest receives from 0.
- Everyone can do it

Parallel Code 2

```
1. /* Add up the integrals calculated by each process */
2. if (my_rank != 0)
3.     MPI_Send(&local_int, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
4.
5. else {
6.
7.     total_int = local_int;
8.     for (source = 1; source < comm_sz; source++) {
9.         MPI_Recv(&local_int, 1, MPI_DOUBLE, source, 0,
10.                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);
11.         total_int += local_int;
12.     }
13. }
```

- All process send local_int to process 0
- Process 0 receives and adds up all local_ints.

Get Input

```
1. void Get_input(int my_rank, int comm_sz, double* a_p, double* b_p,
2.                 int* n_p) {
3.                 int dest;
4.
5.                 if (my_rank == 0) {
6.                     printf("Enter a, b, and n\n\n");
7.                     scanf_s("%lf %lf %d", a_p, b_p, n_p);
8.                     for (dest = 1; dest < comm_sz; dest++) {
9.                         MPI_Send(a_p, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
10.                        MPI_Send(b_p, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
11.                        MPI_Send(n_p, 1, MPI_INT, dest, 0, MPI_COMM_WORLD);
12.                    }
13.                } else { /* my_rank != 0 */
14.                    MPI_Recv(a_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD,
15.                             MPI_STATUS_IGNORE);
16.                    MPI_Recv(b_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD,
17.                             MPI_STATUS_IGNORE);
18.                    MPI_Recv(n_p, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
19.                             MPI_STATUS_IGNORE);
20.                }
21.            } /* Get_input */
```

MPI_Reduce

```
1. int MPIAPI MPI_Reduce(  
2. void *sendbuf,  
3. void *recvbuf,  
4. int count,  
5. MPI_Datatype datatype,  
6. MPI_Op op,  
7. int root,  
8. MPI_Comm comm  
9. );
```

- The rank of the receiving process

MPI_Op

```
1.     typedef enum _MPI_Op {
2.         MPI_OP_NULL    = 0x18000000,
3.         MPI_MAX        = 0x58000001,
4.         MPI_MIN        = 0x58000003,
5.         MPI_SUM        = 0x58000003,
6.         MPI_PROD       = 0x58000004,
7.         MPI_LAND       = 0x58000005,
8.         MPI_BAND       = 0x58000006,
9.         MPI_LOR        = 0x58000007,
10.        MPI_BOR        = 0x58000008,
11.        MPI_LXOR       = 0x58000009,
12.        MPI_BXOR       = 0x5800000a,
13.        MPI_MINLOC     = 0x5800000b,
14.        MPI_MAXLOC     = 0x5800000c,
15.        MPI_REPLACE     = 0x5800000d
16.    } MPI_Op;
```

Parallel Code 2

```
1. /* Add up the integrals calculated by each process */
2. if (my_rank != 0)
3.     MPI_Send(&local_int, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
4.
5. else {
6.
7.     total_int = local_int;
8.     for (source = 1; source < comm_sz; source++) {
9.         MPI_Recv(&local_int, 1, MPI_DOUBLE, source, 0,
10.                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);
11.         total_int += local_int;
12.     }
13. }
```

- All process send local_int to process 0
- Process 0 receives and adds up all local_ints.

Parallel Code 2 using reduce

```
1. /* Add up the integrals calculated by each process */
2. MPI_Reduce(
3.     &local_int,
4.     &total_int,
5.     1,
6.     MPI_DOUBLE,
7.     MPI_SUM,
8.     0,
9.     MPI_COMM_WORLD);
```