



# Week 8

Convolution

1



Road map

- Intro to Convolution
- Sequential program
- OMP or MPI version
- CUDA Intro and Convolution in CUDA

2

## Convolution



Convolution applies a **filter** or **mask** or **kernel\*** on each element of the input array to obtain a new value, which is a **weighted sum** of a set of **neighboring** input elements

\* The term “kernel” may create confusion in the context of GPUs (recall a CUDA/GPU kernel is a function executed by a GPU)

3

## Convolution and its Applications



- **Convolution** is a widely-used operation in signal processing (1D), image processing(2D), video processing(3D), and computer vision
  - Smoothing, sharpening, or blurring an image
  - Finding edges in an image
  - Removing noise, etc.
- Applications in machine learning and artificial intelligence
  - **Convolutional Neural Networks** (CNN or ConvNets)

\* The term “kernel” may create confusion in the context of GPUs (recall a CUDA/GPU kernel is a function executed by a GPU)

4

Mathematically formulation (Discrete Structures...)



$$y_i = \sum_{j=-r}^r f_{i+j} \times x_i$$

- $x_i \in [x_0, x_1, \dots, x_{n-1}]$  //input array
- $y_i \in [y_0, y_1, \dots, y_{n-1}]$  //output array
- $f_i \in [f_0, f_1, \dots, f_{2r}]$  //filter array


5

Some observations




- The size of the filter is an odd number ( $2r + 1$ ),
  - the weighted sum calculation is symmetric around the element that is being calculated.
- That is, the weighted sum involves  $r$  input elements on each side of the position that is being calculated,
- Which is the reason  $r$  is referred to as the radius of the filter.


6



7



8

**FIGURE 7.3**

A 1D convolution boundary condition.

9

### Ghost Elements



- Calculation of **output elements near the boundaries** (beginning and end) of the input array need to deal with “ghost” elements
  - Different policies
  - (0, replicates of boundary values, etc.)

10




FIGURE 7.4  
A 2D convolution example.

11




FIGURE 7.5  
A 2D convolution boundary condition.

12

## Storing the Mask in Constant Memory



- We can store the mask in **constant memory**
  - The mask is small
  - It is constant
  - It is accessed by all threads
- Constant memory is **cached inside each CPU**, and it is particularly fast.

13

## Storing the Mask in Constant Memory



- Declare the mask as a global variable


```
#define MASK_WIDTH 5  
const float M[MASK_WIDTH];
```

14

## Recall: Data Reuse: Tiling




- To take advantage of data reuse, we divide the input into tiles that can be loaded into shared memory



15

15

## Loading Tiles into Shared Memory



**Challenge:** Input and output tiles have different dimensions  
 $(\text{input tile dimension} = \text{output tile dimension} + 2 \times \text{mask radius})$

## Sequential Code: Initializing....



```

1. void convolve(RGBQUAD** image, int width, int height) {
2.     // Allocate memory for the 2D array
3.     RGBQUAD* t_image = (RGBQUAD*)malloc(height * width * sizeof(RGBQUAD*));
4.     RGBQUAD** resultImage = (RGBQUAD**)malloc(height * sizeof(RGBQUAD**));
5.     if (t_image == NULL || resultImage == NULL) {
6.         printf("Unable to allocate memory.\n");    return;}
7.     for (int i = 0; i < height; i++) {
8.         resultImage[i] = &t_image[i*width];
9.     }
10.
11.    // Set kernel Radius
12.    const int kernelRadius = kernelSize / 2;

```

17

## Sequential Code: actual calculations



```

1. // Apply convolution filter
2. for (int y = 0; y < height; y++) {
3.     for (int x = 0; x < width; x++) {
4.         int sumR = 0, sumG = 0, sumB = 0;
5.         for (int ky = -kernelRadius; ky <= kernelRadius; ky++) {
6.             for (int kx = -kernelRadius; kx <= kernelRadius; kx++) {
7.                 int pixelX = x + kx;
8.                 int pixelY = y + ky;
9.                 // Handle boundary conditions (ghost values assumed to be zero)
10.                if (pixelX >= 0 && pixelX < width && pixelY >= 0 && pixelY < height) {
11.                    sumR += image[pixelY][pixelX].rgbRed * kernel5x5[ky + kernelRadius][kx + kernelRadius];
12.                    sumG += image[pixelY][pixelX].rgbGreen * kernel5x5[ky + kernelRadius][kx + kernelRadius];
13.                    sumB += image[pixelY][pixelX].rgbBlue * kernel5x5[ky + kernelRadius][kx + kernelRadius];
14.                } // end if
15.            } // for
16.        } // for
17.    ...
18.    ...

```

18

## Sequential Code: Updating the result image



```

1. // Apply convolution filter
2. for (int y = 0; y < height; y++) {
3.     for (int x = 0; x < width; x++) {
4.         ...
5.         ...
6.         resultImage[y][x].rgbRed = sumR/256;
7.         resultImage[y][x].rgbGreen = sumG/256;
8.         resultImage[y][x].rgbBlue = sumB/256;
9.         // Ensure RGB values are within valid range
10.        resultImage[y][x].rgbRed = (resultImage[y][x].rgbRed > 255) ? 255 :
11.            (resultImage[y][x].rgbRed < 0) ? 0 : resultImage[y][x].rgbRed;
12.        resultImage[y][x].rgbGreen = (resultImage[y][x].rgbGreen > 255) ? 255 :
13.            (resultImage[y][x].rgbGreen < 0) ? 0 : resultImage[y][x].rgbGreen;
14.        resultImage[y][x].rgbBlue = (resultImage[y][x].rgbBlue > 255) ? 255 :
15.            (resultImage[y][x].rgbBlue < 0) ? 0 : resultImage[y][x].rgbBlue;
16.    }
17. }

```

19

## Sequential Code: updataing input image



```

1. // Copy the result image back to the original image
2. for (int y = 0; y < height; y++) {
3.     for (int x = 0; x < width; x++) {
4.         image[y][x] = resultImage[y][x];
5.     }
6. }
7.
8. free(resultImage);
9. free(t_image);
10. printf("Convolution Applied successfully.\n");
}

```

20