Cuda Thread Divergence

1

Minimizing thread divergence

- Warp – Set of threads that execute the same instruction at a time.
- SIMD – Single Instruction, Multiple Data
- SIMT – Singe instruction, Multiple Threads

2

```
1.    Blah;
2.    Blah;
       …
1.    If ( . . .)
2.    {
3.     //then do smth
4.    }
5.    Else
6.    {
7.    //else do smth
8.    }
9.    Blah;
10.   Blah;
```

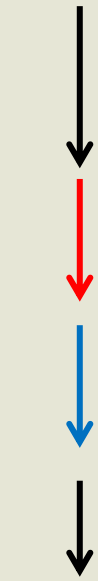

3

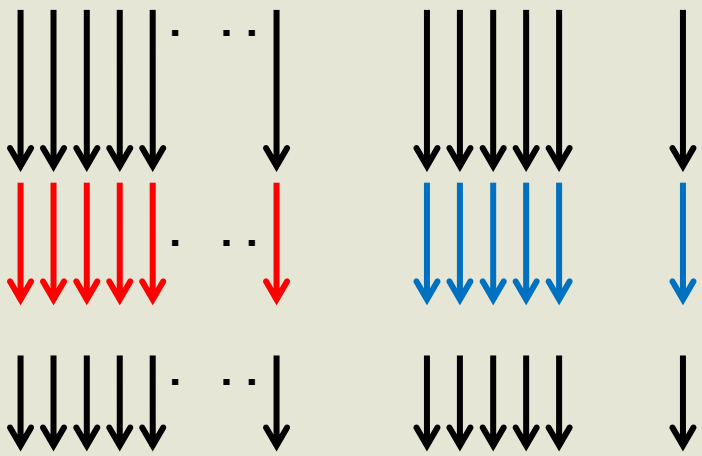

4

2-Way branch divergence

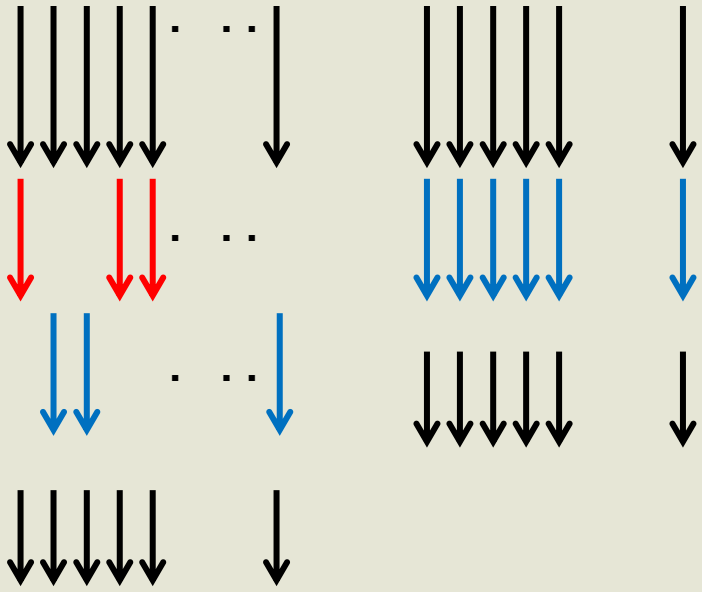

5

What is the max branch divergence penalty for a cuda thread block with 1024 threads?

_____ x slowdown

6

What is the max branch divergence penalty for a cuda thread block with 1024 threads?

32 x slowdown

Max 32-way branch divergence (warp size)

7

```
1.  Switch (expr) {
2.       case 1: . . . Break;
3.       case 2: . . . Break;
4.       case 3: . . . Break;
5.       …
6.       case 32: . . . Break;
7.  }
```

8

Survey

1. `Switch (threadidx.x %32) {case 0..31}`
2. `Kernel <<< 1, 1024 >>>();` _____

3. `Switch (threadidx.x % 64) {case 0..63}`
4. `Kernel <<< 1, 1024 >>>();` _____

5. `Switch (threadidx.y) {case 0..31}`
6. `Kernel <<< 1, 64 x 16 >>>();` _____

7. `Switch (threadidx.y) {case 0..31}`
8. `Kernel <<< 1, 16 x 16 >>>();` _____

- What will be the slowdown for each of the following expressions in switch statements

9

Survey

1. `Switch (threadidx.x % 2) {case 0..31}`
2. `Kernel <<< 1, 1024 >>>();` _____

3. `Switch (threadidx.x / 32) {case 0..31}`
4. `Kernel <<< 1, 1024 >>>();` _____

5. `Switch (threadidx.x / 8) {case 0..63}`
6. `Kernel <<< 1, 1024 >>>();` _____

- What will be the slowdown for each of the following expressions in switch statements

10

Branch Divergence in real world

- Assume a 1024 x 1024 image.
- Requiring Special handling of pixels on the boundary

11

Maximum branch divergence of any warp? _____ - way

```
1.  __global__
2.  If(threadIdx.x == 0 ||
       threadxIdx.x == 1024
       threadIdx.y==0 ||
       threadIdx.y==1024){
1.     //deal with boundary cond.
2.  }else {
3.     //do smth
4.  }
```

12

Branch divergence in convolution

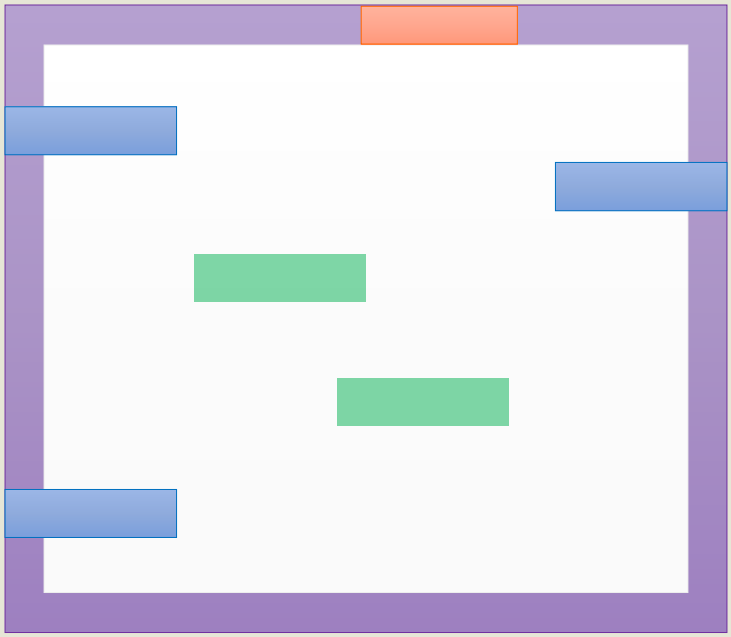

13

Branch divergence in real life

- Be aware of branch divergence
- Don’t panic if there are if statements
- No real strategy in reducing branch divergence

14

Major guidelines

- Avoid code with too many branches
- Be aware of large imbalance in thread workloads
  - For loops with variable terminating statements.

15

TILED MATRIX OPERATIONS

16

N X N MATRIX

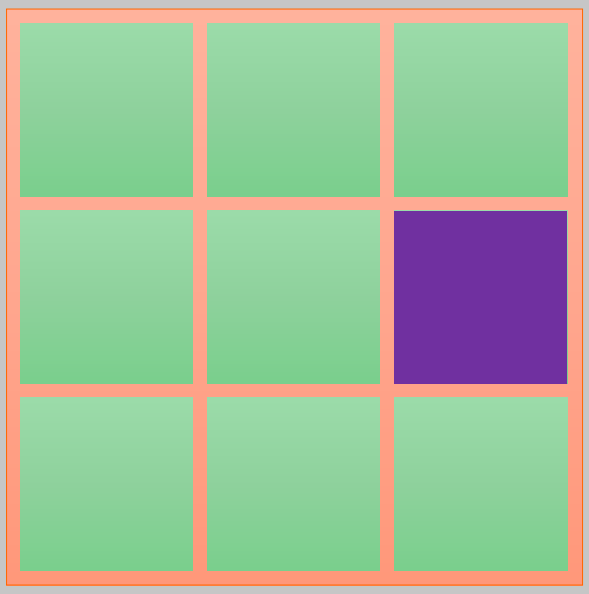

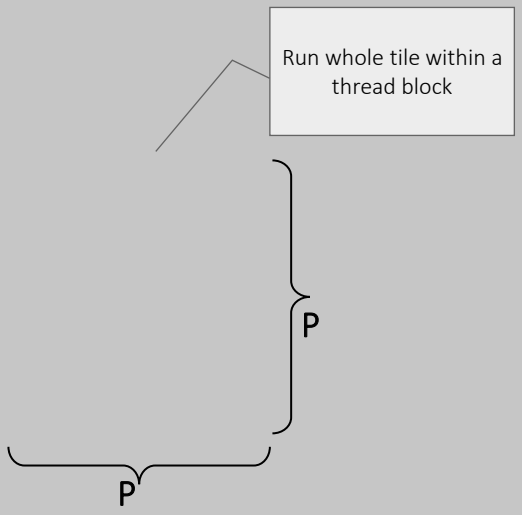

Tile P x P So N^2/P^2 tiles overall

17

P^2 THREADS TO COMPUTE 1 P X P TILE

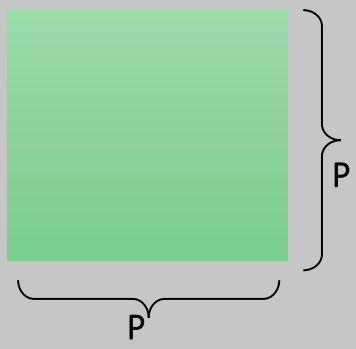

- Good: P^2 parallel ops
- Bad: Must share parameters btw threads
- How many threads must get the parameters for
  - Each source element?___
  - Each dest. Element?____

18

P THREADS TO COMPUTE 1 P X P TILE

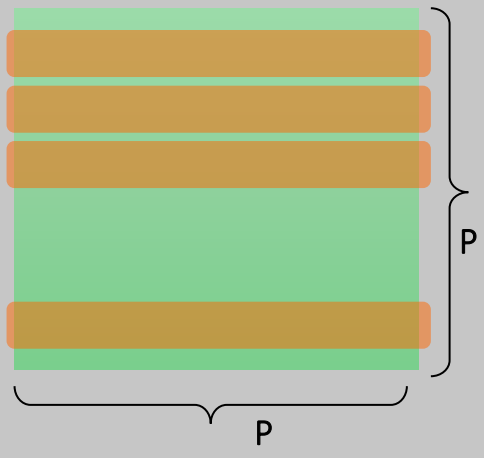

Is parallelization

increased? _____

decreased?_____

No change?_____

19

P THREADS TO COMPUTE 1 P X P TILE

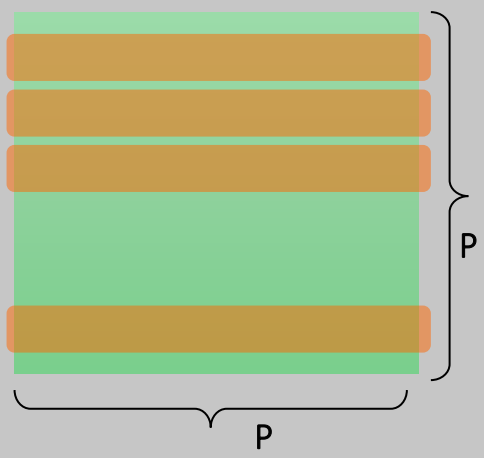

- fewer threads
- More work per thread
- Communication
  - among threads vs within a thread

20